

# CMPU4003 Advanced Databases

## Working with JSONB in PostgreSQL

**Goal:** Learn practical patterns for storing, querying, updating, and indexing semi-structured academic data (students, subjects, grades, institutions) using jsonb in PostgreSQL.

---

### 1. Setup

- If you have a PostgreSQL installation from last year this should be sufficient and you do not need to setup a new one.
- If you need to setup PostgreSQL there are a number of options:
  - Setup a Docker installation/Dev Container
    - Follow the instructions in Option 1 Docker from command line.txt
    - OR
    - Follow the instructions in Option 2 Docker in an IDE.txt
    - OR
    - Follow the instructions in Option 3 Dev Container in VS Code.txt
  - Use Supabase.com
    - Follow the instructions in Option 4 Supabase.txt
- Once you have PostgreSQL setup create a schema for this lab and set the search path so that Postgres will use this schema:

```
CREATE SCHEMA IF NOT EXISTS jsonb_lab;  
SET search_path = jsonb_lab;
```

### 2. Practical Tasks

Universities often record structured data (student IDs, subject codes) and semi-structured data (metadata about assessments, remarks, evolving grading rubrics). jsonb helps store flexible information without altering schemas constantly. We are going to model students, enrollments and grades as JSON B.

## 2.1 Create and populate the tables in the jsonb\_lab schema.

You are going to create two tables students and enrollments.

In students you will have a column Profiles which will store personal metadata (age, major, languages, sports, exchange) of type JSONB.

In enrollments you will have a column Grades which will store different assessment structures (simple scores, arrays of assignments, nested project details) of type JSONB.

### # Drop and Create Tables

```
DROP TABLE IF EXISTS students CASCADE;
CREATE TABLE students (
  student_id SERIAL PRIMARY KEY,
  name TEXT NOT NULL,
  university TEXT NOT NULL,
  profile JSONB NOT NULL DEFAULT '{}'::jsonb
);

DROP TABLE IF EXISTS enrollments;
CREATE TABLE enrollments (
  enrollment_id BIGSERIAL PRIMARY KEY,
  student_id INT REFERENCES students(student_id),
  subject_code TEXT NOT NULL,
  year INT NOT NULL,
  semester INT NOT NULL,
  grades JSONB NOT NULL -- stores assignments, exams, comments
);
```

### -- Insert Sample Data

```
INSERT INTO students (name, university, profile) VALUES
('Alice Johnson', 'TU Dublin', '{"age":22,"major":"CS","languages":["en","fr"]}'),
('Brian Smith', 'UCD', '{"age":24,"major":"Math","sports":["football"]}'),
('Chloe Lee', 'Trinity', '{"age":21,"major":"Engineering","exchange":true}');

INSERT INTO enrollments (student_id, subject_code, year, semester, grades) VA
LUES
(1, 'DB4003', 2023, 1, '{"midterm":78,"final":85,"remarks":"Good progress"}')
,
(1, 'ML4001', 2023, 2, '{"assignments":[{"name":"A1","mark":40},{"name":"A2","mark":45}], "final":82}'),
(2, 'DB4003', 2023, 1, '{"midterm":65,"final":70,"remarks":"Needs work"}'),
(3, 'CS4090', 2023, 2, '{"project":{"title":"IoT","mark":88},"oral_exam":90}')
);
```

Explanation:

Note that profile is flexible: Each student has a slightly different structure (some have languages, some sports, some exchange).

<p>Alice Johnson</p> <pre>{   "age": 22,   "major": "CS",   "languages": ["en", "fr"] }</pre> <p>Keys: age (number), major (string), languages (array of strings).</p>	<p>Brian Smith</p> <pre>{   "age": 24,   "major": "Math",   "sports": ["football"] }</pre> <p>Keys: age, major, sports (array of strings).</p>	<p>Chloe Lee</p> <pre>{   "age": 21,   "major": "Engineering",   "exchange": true }</pre> <p>Keys: age, major, exchange (boolean).</p>
--	--	--

## 2.2 JSONB Basics

->> means **extract a JSON field as text**.

-> means **extract a JSON field as JSON**.

#> lets you navigate deeper using a **path array**.

the ? operator means: **“Does this JSON object contain the given key?”**

:: is the **type cast operator**.

Try these queries:

*-- Extract fields*

```
SELECT name,
       profile->>'major' AS major,
       (profile->>'age')::int AS age
FROM students;
```

*-- Extract fields as Json and text to illustrate the operator*

*--no different in visual editor will matter in applications.*

```
SELECT name,
       profile->'major' AS major_json,
       profile->>'major' AS major_text
FROM students;
```

```

-- Extract fields casting age to be type integer
SELECT name,
       profile->>'major' AS major,
       (profile->>'age')::int AS age
FROM students;

-- Does profile contain exchange info?
SELECT name FROM students WHERE profile ? 'exchange';

-- Subjects with final mark >= 80 where final is cast as an integer
SELECT subject_code, (grades->>'final')::int AS final
FROM enrollments
WHERE (grades->>'final')::int >= 80;

-- Suppose we are looking for students with languages as part of their profile.
-- We know Alice has a profile with {"languages":["en","fr"]}
-- profile#>'{languages}' → extracts the whole array as JSON.
-- profile#>>'{languages,0}' → navigates into the array (0 = first element) and returns text.
SELECT name,
       profile#>'{languages}' AS langs_json,
       profile#>>'{languages,0}' AS first_lang
FROM students
WHERE profile ? 'languages';

```

## 2.3 Working with Arrays and Nested Objects

```

-- Expand assignments into rows to get the grades for each assignment for subject ML4001
SELECT e.subject_code, a->>'name' AS assignment, (a->>'mark')::int AS mark
FROM enrollments e
CROSS JOIN LATERAL jsonb_array_elements(e.grades->'assignments') a
WHERE subject_code = 'ML4001';

```

Explanation:

**e.grades->'assignments'**

→ gets the value of the "assignments" key from the grades JSON.

**jsonb\_array\_elements(...)** a

→ takes that JSON array and unnests it into multiple rows.

- Row 1: {"name":"A1","mark":40}
- Row 2: {"name":"A2","mark":45}

**CROSS JOIN LATERAL**

→ means: for each row in enrollments, run this function and join the results.

- Without **LATERAL**, you can't pass values from the left table (e.grades) into the function.

a->>'name'

→ extracts "A1" / "A2" as text. (a->>'mark')::int

→ extracts "40" / "45" as text, then casts to integer.

```
-- Extract project marks
SELECT subject_code, grades#>>'{project,title}' AS project_title,
       (grades#>>'{project,mark}'):int AS mark
FROM enrollments
WHERE grades ? 'project';
```

Explanation:

**WHERE grades ? 'project'**

→ ensures we only look at rows where the grades JSON has a "project" key.

**grades#>>'{project,title}'**

→ use #>> (**path operator**) to navigate nested JSON:

- Go into "project"
- Extract "title" as text

**(grades#>>'{project,mark}'):int**

→ go into "project" → extract "mark" as text → cast to integer.

## 2.4 JSONPath Queries

A **JSON path** is like a *query language* (a bit like XPath for XML) that lets you navigate inside a JSON document.

- Think of a JSON document as a tree of objects and arrays.
- A JSON path is a string (starting with \$) that says “**go here**” inside that tree.

Examples:

- \$ → the root of the JSON document
- \$.assignments → the assignments field
- \$.assignments[\*].mark → all the mark values inside the assignments array
- \$.\* → all the fields at the root, whatever their names

```
-- Students with any grade (for anything assignment, final etc) >= 85
SELECT enrollment_id, subject_code
FROM enrollments
WHERE jsonb_path_exists(grades, '$.* ? (@ >= 85)');
```

Explanation:

**jsonb\_path\_exists(grades, ...)**

Checks if the JSON path finds at least one match inside the grades JSON. Returns true or false.

Path: '\$.\* ? (@ >= 85)'

- \$ = root of the JSON document (grades).
- .\* = all keys at the root (like "midterm", "final", "remarks", "assignments", "project", etc.).
- ? (@ >= 85) = filter: return only the values >= 85.

```
-- For subject ML4001, show every assignment mark stored in the grades JSON."
```

```
SELECT jsonb_path_query(grades, '$.assignments[*].mark')
FROM enrollments
WHERE subject_code = 'ML4001';
```

Explanation:

**jsonb\_path\_query(grades, ...)**

Extracts the values that match the given path.

Path: \$.assignments[\*].mark

- \$ = root.
- .assignments = go into the assignments key.
- [\*] = all elements of the array.
- .mark = take the mark field of each.

## 2.5 Aggregations

```
-- Average final grade per university
SELECT s.university, AVG((e.grades->>'final')::int)
FROM students s
JOIN enrollments e ON s.student_id = e.student_id
WHERE e.grades ? 'final'
GROUP BY s.university;
```

Explanation:

**JOIN**: links each enrollment (e) with the corresponding student (s) so we can access both the grades and the student's university.

**e.grades ? 'final'**: only keep rows where the JSON grades has a key "final".

**e.grades->>'final'**: extract the final grade from JSON as text.

**::int**: cast it to an integer so math can be done.

**AVG(...)**: compute the average per group.

```
-- Best student per subject
```

```

SELECT subject_code, student_id, MAX((grades->>'final')::int) AS best
FROM enrollments
WHERE grades ? 'final'
GROUP BY subject_code, student_id;

```

Explanation:

`grades ? 'final'`: filter to rows where "final" exists.

`grades->>'final'`: extract the final grade (text).

`::int`: cast to integer.

`MAX(...)`: compute the maximum final grade.

`GROUP BY subject_code, student_id`: groups by subject and student.

## 2.6 Updates

*-- Add exchange flag for all UCD students*

```

UPDATE students
SET profile = profile || '{"exchange": false}'
WHERE university = 'UCD';

```

Explanation:

`profile` is a JSONB column.

`||` is the **concatenation / merge operator** for JSONB.

It merges the existing profile object with `{"exchange": false}`.

- If exchange already exists, it will be **overwritten** with false.
- If not, the key is added.

*-- Update a nested grade*

*-- For student 2 in DB4003, set their final grade inside the grades JSON to 90.*

```

UPDATE enrollments
SET grades = jsonb_set(grades, '{final}', '90')
WHERE subject_code = 'DB4003' AND student_id = 2;

```

Explanation:

- `jsonb_set(target, path, new_value)` replaces or inserts a value at the given path.
- `grades` is the JSONB column.
- `'{final}'` is the path (an array with one key, "final").
- `'90'` is the new value (a JSON number here, since no quotes inside).
- Only applies to the enrollment where `subject = DB4003` and `student_id = 2`.

*-- Remove remarks for enrollments in subject DB4003*

```
UPDATE enrollments
SET grades = grades - 'remarks'
WHERE subject_code = 'DB4003';
```

Explanation:

- 'key' removes a key from a JSONB object.

This removes the "remarks" field from the grades JSON.

Only for enrollments in DB4003.

### Exercise:

1. Add a new key ECTS = 5 to all DB4003 enrollments.
2. Remove the key midterm where present.

## 2.8 Constraints

*-- Ensure grades are JSON objects*

```
ALTER TABLE enrollments
ADD CONSTRAINT grades_is_object CHECK (jsonb_typeof(grades) = 'object');
```

*-- Ensure final mark between 0-100*

```
ALTER TABLE enrollments
ADD CONSTRAINT final_between CHECK ((grades ? 'final') IS NOT TRUE OR ((grades->>'final')::int BETWEEN 0 AND 100));
```

Explanation: Checking that either it doesn't exist or that if it does that it is between 0 and 100

## 3. Exercises

1. Create a view subject\_results with student name, subject, year, final grade.  
Create or Replace View....
2. Find top 3 students in CS4090 (by any grade).  
For CS4090, students may have "project.mark" or "oral\_exam". We can take the **\*\*maximum numeric value inside `grades`\*\*** and rank by that.  
Remember use Limit to limit your results
3. Using JSONPath, find students with assignment average > 40.



Join to students

CROSS JOIN LATERAL jsonb\_array\_elements(e.grades->'assignments') AS a

Remember to use Group and Having

4. Add a constraint ensuring that if oral\_exam exists, its value is  $\leq 100$ .